

A cut-cell immersed boundary method in a collocated finite-volume solver

Formulation, implementation in OpenFOAM and verification

September 29, 2026

Abstract

We present a cut-cell immersed boundary method for the incompressible Navier–Stokes equations on a fixed, non body-fitted mesh. The body is described by a level set; each face of the mesh carries the fraction of its area that lies in the fluid and each cell the fraction of its volume. The wall segment of a cut cell is not computed independently but *defined* by the geometric closure of the fluid polyhedron, and this single identity carries the whole boundary treatment: the wall pressure force is absorbed into a face sum, no-penetration follows from discrete continuity, and the difference between a free-slip and a no-slip wall is one diagonal coefficient. No ghost values, no extrapolation into the solid and no coupling between velocity components appear. The method is implemented in OpenFOAM without modifying the mesh: the fractions enter the standard operators explicitly, which keeps the pressure–velocity coupling of `simpleFoam`/`pimpleFoam` and the linear algebra untouched. The implementation is verified against an independent reference implementation of the same discretisation on a cylinder at $Re_D = 1$.

Contents

1	Introduction	2
2	Governing equations	2
3	Geometric description of the body	3
3.1	Level set, fluid fractions and cell classes	3
3.2	The closure identity	3
3.3	Computation of the fractions	4
3.4	Sliver absorption and coupled faces	5
3.5	Derived quantities	5
3.6	Bodies	6
4	Discretisation for a body at rest	7
4.1	Term-by-term derivation	7
4.2	Properties	8
5	Pressure–velocity coupling, forces and flow-rate forcing	9
5.1	Solid cells	9
5.2	SIMPLEC with cut-cell operators	10
5.3	Compatibility and the pressure reference	11
5.4	Algorithm	11
5.5	Force on the body	11
5.6	Constant flow rate in periodic domains	12

6	Implementation in OpenFOAM	12
6.1	Explicit operators on an unmodified mesh	12
6.2	Structure	12
6.3	Input and output	13
7	Verification	13
7.1	Body at rest	13
8	Limitations	14
9	Conclusions	15

1 Introduction

Immersed boundary methods let a body of arbitrary shape be represented on a mesh that does not conform to it. Among them, cut-cell methods are distinguished by being conservative: the fluid part of every cell is a well-defined control volume, and fluxes are accounted for on its actual boundary. Their classical difficulties are the description of the cut geometry, the small-cell problem.

The method described here addresses these points with a minimum of machinery. Its central object is the *closure identity* of a cut cell: the wet parts of its faces and its wall segment form a closed surface, so the wall area vector equals minus the sum of the wet face area vectors. Rather than computing the wall segment from the geometry and then enforcing consistency, the method defines the wall segment through the identity. Every consequence of the boundary condition — pressure force, no-penetration, wall traction — is then expressed with the same face fractions that the flux operators use, and the resulting discrete properties (exactness for uniform flow, no spurious force from a uniform pressure, exact compatibility of the pressure system) hold to round-off by construction.

The document is organised as follows. Section 2 states the equations, Section 3 defines the geometry and how it is computed, Section 4 derives the discretisation for a body at rest and its properties, and Section 5 the pressure–velocity coupling, the forces and the flow-rate forcing. Section 6 describes the implementation in OpenFOAM, Section 7 the verification, and Section 8 the limitations.

2 Governing equations

We solve the incompressible Navier–Stokes equations

$$\frac{\partial \mathbf{u}}{\partial t} + \nabla \cdot (\mathbf{u}\mathbf{u}) - \nabla \cdot (\nu \nabla \mathbf{u}) = -\nabla p + \mathbf{g}, \quad \nabla \cdot \mathbf{u} = 0, \quad (1)$$

with p the kinematic pressure, ν the kinematic viscosity and $\mathbf{g}$ a body force, in a domain Ω containing one or more rigid bodies B_i at rest. On each body surface one of two conditions is imposed: *no slip*, $\mathbf{u} = \mathbf{0}$, or *free slip*, $\mathbf{u} \cdot \mathbf{n} = 0$ together with zero tangential traction, $(\mathbf{I} - \mathbf{n}\mathbf{n}) \cdot (\nabla \mathbf{u} + \nabla \mathbf{u}^\top) \cdot \mathbf{n} = \mathbf{0}$.

The discretisation is a collocated finite-volume method with SIMPLE(C) pressure–velocity coupling in the arrangement of OpenFOAM’s `simpleFoam` (steady) and `pimpleFoam` (transient, implicit Euler with outer iterations). We write V_P for the volume of cell P , $\mathbf{S}_f$ for the area vector of face f pointing out of P , N for the neighbour across f , ϕ_f for the volume flux through f , $\mathbf{d}_f = \mathbf{x}_N - \mathbf{x}_P$ and $\delta_f = 1/|\mathbf{d}_f|$ for the inverse of the centre-to-centre distance. Face values $(\cdot)_f$ are linear interpolates, $\frac{1}{2}[(\cdot)_P + (\cdot)_N]$ on a uniform mesh. Sums $\sum_f$ run over the faces of the cell under consideration.

3 Geometric description of the body

3.1 Level set, fluid fractions and cell classes

Definition 1 (Level set and fractions). *A body is given by a function $\varphi(\mathbf{x})$, negative inside the body and positive in the fluid; only its sign and its zero level are used, and $\varphi = 0$ counts as solid. For a face f and a cell P ,*

$$\theta_f = \frac{|\{\mathbf{x} \in f : \varphi > 0\}|}{|\mathbf{S}_f|} \in [0, 1], \quad \alpha_P = \frac{|\{\mathbf{x} \in P : \varphi > 0\}|}{V_P} \in [0, 1] \quad (2)$$

are the fluid area and volume fractions. Several bodies are combined through the pointwise minimum of their level sets, which represents their union exactly.

Cells fall into three classes (Fig. 1a): *fluid* ($\alpha_P = 1$), *solid* ($\alpha_P = 0$) and *cut* ($0 < \alpha_P < 1$). A fourth set, overlapping the others, collects the *wall cells*: every cell with $\alpha_P > 0$ whose wall area (defined below) exceeds $10^{-10}V_P^{2/3}$. It contains all cut cells and, in addition, the fluid cells one of whose faces has been closed by a solid neighbour (Section 3.4). Wall cells are the only cells in which a wall term appears.

The fluid part Ω_P of a cut cell is a polyhedron bounded by the wet parts of the cell faces, of area vectors $\theta_f \mathbf{S}_f$, and by the wall segment, of area A_{wall} and outward unit normal $\mathbf{n}_w$ (Fig. 1b). Its volume is $\alpha_P V_P$.




Figure 1: Geometry of the cut cells, computed with the algorithm of Section 3.3 for a cylinder of radius 1 on a mesh of spacing $h = 0.2$. (a) Cell classes. (b) One cut cell: corner signs (open squares fluid, filled squares solid), the bisection iterates (red ticks) converging to the edge crossings (red dots), the wet parts of the faces (green), the chord joining the crossings (dashed), and the sub-sampling lattice that gives α_P (blue points in the fluid, grey in the body; 8×8 drawn, 32×32 used). (c) The closure polygon of the same cell: the wet face area vectors, laid head to tail, are closed by the wall area vector $\mathbf{S}_w$ (Proposition 1).

3.2 The closure identity

Proposition 1 (Closure). *For every cell, $\sum_f \theta_f \mathbf{S}_f + \mathbf{n}_w A_{\text{wall}} = \mathbf{0}$.*

Proof. Apply the divergence theorem to a constant vector field $\mathbf{a}$ over Ω_P : $0 = \int_{\Omega_P} \nabla \cdot \mathbf{a} \, dV = \mathbf{a} \cdot (\sum_f \theta_f \mathbf{S}_f + \mathbf{n}_w A_{\text{wall}})$ for every $\mathbf{a}$. For a planar wall segment $\mathbf{n}_w A_{\text{wall}}$ is its area vector; for a curved one it is the integral $\int_{\text{wall}} \mathbf{n} \, dS$, and the identity holds unchanged. $\square$

We write

$$\mathbf{S}_w \equiv - \sum_f \theta_f \mathbf{S}_f = \mathbf{n}_w A_{\text{wall}}, \quad A_{\text{wall}} = |\mathbf{S}_w|, \quad (3)$$

for the wall area vector, pointing from the fluid into the body. In the implementation $\mathbf{S}_w$ is never computed from the level set: it is *defined* by (3) from the same θ_f that weight every flux operator, so the identity holds to round-off however the wall is inclined or curved. Figure 1c shows the construction. Two consequences are used throughout. First, every property that follows from the divergence theorem for a closed surface holds discretely: a constant field has zero net flux, a uniform pressure exerts no force. Second, the geometric errors of the fractions (the chord instead of the arc on a face, the quadrature error of α_P) change *where* the discrete wall is, but never make the discrete control volume leak.

In a cell whose faces are all fully open or fully closed, (3) also gives a non-zero $\mathbf{S}_w$ as soon as one face is closed: for a fluid cell next to a solid neighbour, $\mathbf{S}_w$ is the area vector of the closed face, i.e. the cell has a staircase wall. The two cases need no distinction in any operator.

3.3 Computation of the fractions

The fractions are computed once, when the solver starts, as follows.

1. *Evaluation.* The union level set is evaluated at the mesh points, the face centres and the cell centres.
2. *Edge crossings.* On every edge whose end points have different fluid status, the crossing is located by bisection of the exact level set, $n_{\text{bis}} = 40$ steps by default. The edge is always traversed in a canonical direction, from the lexicographically smaller end point A to the larger one B : with bracket $[\lambda_{\text{lo}}, \lambda_{\text{hi}}] \subset [0, 1]$, the midpoint $\mathbf{x}_m = A + \frac{1}{2}(\lambda_{\text{lo}} + \lambda_{\text{hi}})(B - A)$ replaces the end of the bracket that has the same fluid status, and the crossing is the final midpoint. Its error is at most $2^{-41} \approx 4.5 \times 10^{-13}$ of the edge length. The canonical direction makes the two copies of an edge on a processor boundary, which have identical coordinates, obtain bit-identical crossings; across a translational periodic boundary the lexicographic order is preserved, so the two sides agree whenever the level set is itself periodic. All cut edges are bisected together, one level-set call per iteration, which matters for bodies whose level set is expensive (triangulated surfaces).
3. *Face fractions.* A face whose vertices are all fluid has $\theta_f = 1$, one whose vertices are all solid $\theta_f = 0$. Otherwise the vertices are walked in order; each fluid vertex is appended to a polygon, and after every vertex whose status differs from the next the crossing of that edge is appended. The polygon area $|\sum_i \frac{1}{2}(\mathbf{p}_i - \mathbf{p}_0) \times (\mathbf{p}_{i+1} - \mathbf{p}_0)|$ divided by $|\mathbf{S}_f|$ is θ_f . A face with alternating signs (a saddle) is resolved as *fluid-connected*: the walk yields one polygon joining both fluid corners. The wet part of the face is thus approximated by the polygon whose edges are chords of the wall's trace on the face, a second-order approximation in the mesh spacing.
4. *Cell fractions.* A cell is a candidate for sub-sampling when the level set at its centre, its face centres and its vertices does not have one sign; the others have $\alpha_P = 1$ or 0. For a candidate, α_P is the fraction of fluid points in a lattice of n_s points per direction ($n_s = 32$ by default) placed at the sub-cell centres of the cell's bounding box,

$$\alpha_P = \frac{\#\{k : \varphi(\mathbf{x}_k) > 0\}}{\#\{k : \mathbf{x}_k \in P\}}, \quad (4)$$

where samples outside the cell are discarded (the point-in-cell test is skipped for axis-aligned hexahedra, whose bounding box is the cell). In a direction that is not solved (the empty direction of a two-dimensional case) a single sample at the cell centre is taken, so

a two-dimensional case samples n_s^2 points and a three-dimensional one n_s^3 . This midpoint quadrature has an error of order $1/n_s$ per cell for a generic cut (about 10^{-3} for $n_s = 32$), which is small compared with the discretisation error.

5. *Sliver absorption and coupled faces*, Section 3.4.

6. *Derived quantities*, Section 3.5.

Two features smaller than a cell are not resolved: a body that crosses a face without crossing any of its edges leaves θ_f unchanged, and one that enters a cell without reaching its centre, a face centre or a vertex leaves $\alpha_P = 1$. Note also that α_P and θ_f come from different algorithms and are therefore not exactly consistent (the polyhedron implied by the θ_f does not have exactly the volume $\alpha_P V_P$); nothing in the method requires them to be, because α_P enters only as a volume weight and $\mathbf{S}_w$ is built from θ_f alone.

3.4 Sliver absorption and coupled faces

A cut cell with a tiny fluid volume but open faces is the source of the small-cell problem. The fractions are therefore post-processed in two steps (Fig. 2).

Absorption. Every cell with $\alpha_P < \alpha_{\min}$ ($\alpha_{\min} = 0.02$ by default) is absorbed into the body: $\alpha_P := 0$ and *all its faces are closed*, $\theta_f := 0$. The same rule closes the faces of every cell that has $\alpha_P = 0$ from the start, so a solid cell never has an open face. Closing the faces is essential: an open face between an absorbed cell and a fluid neighbour would let the neighbour exchange flux with a cell that carries no equation, and the closure identity would then place the neighbour's wall on a face that is not a wall. After absorption the neighbours see the closed face through (3) as part of their wall (Fig. 2b), so the treatment is automatically consistent. The absorption displaces the boundary by at most $O(\alpha_{\min} h)$ and removes the fluid volume $\alpha_P V_P < \alpha_{\min} V_P$ of each absorbed cell. The pass is not iterated: a neighbour that loses faces keeps its α .

Coupled faces. The two copies of a face on a periodic or processor boundary must carry the same fraction. After absorption, θ_f is synchronised across coupled faces with the minimum of the two copies,

$$\theta_f := \min(\theta_f, \theta_{f'}), \quad (5)$$

which also propagates the closing of an absorbed cell across the boundary. The order matters (Fig. 2c): synchronising first and absorbing second would close the face on the side of the absorbed cell only, and its partner cell would keep an open face into a cell without equation. A difference larger than 10^{-10} between the two copies before synchronisation is reported as a warning (a body that is not periodic although the mesh is).

3.5 Derived quantities

From the fractions the following fields are formed.

- The wall vector $\mathbf{S}_w$ by (3), accumulated face by face: each internal face subtracts $\theta_f \mathbf{S}_f$ from its owner and adds it to its neighbour, each boundary face subtracts it from its cell. Boundary faces of the domain count with their own θ_f , so a body that cuts an inlet or a wall patch still has a closed control volume. The front and back faces of a two-dimensional case do not enter (they cancel).
- The wall area $A_{\text{wall}} = |\mathbf{S}_w|$ and the wall normal $\mathbf{n}_w = \mathbf{S}_w / A_{\text{wall}}$.
- The *wall distance* of the cell value from the wall,

$$d_{\text{wall}} = \frac{\alpha_P V_P}{2 A_{\text{wall}}}. \quad (6)$$




Figure 2: Sliver absorption, with real fractions for a cylinder of radius 1 on a mesh of spacing 0.2. Closed faces are black lines; red arrows are the wall vectors $\mathbf{S}_w$ of the wall cells (scaled by 0.8), drawn from the fluid centroid. (a) A sliver with $\alpha_P = 0.018$. (b) After absorption its faces are closed, and the neighbour above it becomes a wall cell whose $\mathbf{S}_w$ now includes the closed face. (c) On a periodic or processor boundary the coupled face must be closed on both sides, which requires absorbing *before* synchronising with the minimum.

It is exact for a fluid slab: if the fluid part of a cell is a layer of thickness δ on a wall of area A parallel to a face, then $\alpha_P V_P = A\delta$, $A_{\text{wall}} = A$, and $d_{\text{wall}} = \delta/2$ is the distance from the centroid of the layer to the wall. For a general cut it is an estimate of that distance from two integral quantities that the closure identity delivers exactly; Fig. 3 compares it with the true centroid distance. For a fluid cell with one closed face on a Cartesian mesh it gives $h/2$, the exact distance of the cell centre from the face.

- The *no-slip (Dirichlet) coefficient*

$$c_{\text{ns}} = \frac{A_{\text{wall}}}{d_{\text{wall}} V_P} = \frac{2A_{\text{wall}}^2}{\alpha_P V_P^2} \quad [\text{m}^{-2}] \quad (7)$$

in wall cells of no-slip bodies (and of wall cells without a body), zero elsewhere. Multiplied by a diffusivity it is the wall exchange coefficient per unit volume.

- The *blanking coefficient* $b_P = \sum_f |\mathbf{S}_f| \delta_f / V_P$ in solid cells, zero elsewhere: the viscous diagonal per unit volume and diffusivity of an unobstructed cell.
- The body index of every cut, solid and wall cell (Section 3.6).

3.6 Bodies

Table 1 lists the body descriptions. Every body carries a wall condition (`wallType noSlip` or `freeSlip`) and may be inverted ($\varphi \rightarrow -\varphi$, e.g. fluid inside a cylinder).

Table 1: Body descriptions. $\mathbf{c}$ centre, R radius, $\hat{\mathbf{a}}$ unit axis, $\mathbf{h}$ half-size and $\mathbf{q} = |\mathbf{x} - \mathbf{c}| - \mathbf{h}$ componentwise, $\hat{\mathbf{n}}$ unit normal into the fluid.

type	level set $\varphi(\mathbf{x})$	remark
cylinder	$ \mathbf{r} - (\mathbf{r} \cdot \hat{\mathbf{a}})\hat{\mathbf{a}} - R$, $\mathbf{r} = \mathbf{x} - \mathbf{c}$	infinite
sphere	$ \mathbf{x} - \mathbf{c} - R$	
box	$ \max(\mathbf{q}, \mathbf{0}) + \min(\max_i q_i, 0)$	exact signed distance
plane	$(\mathbf{x} - \mathbf{x}_0) \cdot \hat{\mathbf{n}}$	half space
any closed searchableSurface	$\pm \mathbf{x} - \mathbf{x}_{\text{nearest}} $, $-$ inside	e.g. STL
alphaField	$-$ (fraction field supplied)	see text




Figure 3: Wall distance. (a) In a cut cell the first-order estimate (6) is half the thickness $\alpha V / A_{\text{wall}}$ of the slab with the same volume on the same wall area; the distance d_c of the fluid centroid $\mathbf{x}_c$ from the wall is the quantity it approximates. (b) Ratio of the two over all cut cells of a cylinder of radius 1 for 40 random offsets of a mesh of spacing 0.1: the estimate is within 25% for $\alpha_P \lesssim 0.8$ and overestimates the distance in cells cut near a corner, whose small wall area makes d_{wall} large (and c_{ns} small, so the wall term there is weak but also unimportant).

A closed triangulated surface (any OpenFOAM `searchableSurface` with an inside/outside test, e.g. `triSurfaceMesh` reading an STL file) uses the signed distance to the surface as level set, so that the same bisection applies. A surface without a volume type is rejected. A body given as a fluid-fraction field α is intended for fractions produced by another model. Its cell fractions are the field values (clipped to $[0, 1]$, optionally converted from a solid fraction); point values are obtained by inverse-distance interpolation, $\alpha - \frac{1}{2}$ serves as a point level set, edge crossings are found by linear interpolation of that level set, and the same polygon construction gives θ_f . This route smears the wall over about one cell (Section 8).

The level sets of all bodies with a level set are combined by their minimum before the fractions are computed, so overlapping bodies form an exact union; fraction-field bodies are merged afterwards with the minimum of α and of θ . Every cut or solid cell is attributed to the body with the most negative level set at its centre (for a fraction-field body: every cell whose fraction it lowered), and a wall cell that is not cut inherits the body across its closed face. The attribution selects the wall condition, the wall body to which the cell's force contribution is added. A wall cell without a body is treated as no-slip and its force is added to the first body.

4 Discretisation for a body at rest

Every term of (1) is integrated over the fluid part Ω_P of each cell. The boundary of Ω_P consists of the wet parts of the faces and of the wall segment. The face parts are treated as in any finite-volume method, weighted by θ_f ; the wall segment contributes to each term according to the boundary condition, and this is the only place where the boundary condition enters. Table 2 collects the result, and the subsections derive each entry.

4.1 Term-by-term derivation

Unsteady term. $\int_{\Omega_P} \partial_t \mathbf{u} dV = \alpha_P V_P d\mathbf{u}_P / dt$ for a cell whose fluid volume does not change, discretised with implicit Euler. The cell value $\mathbf{u}_P$ is the average over the fluid part.

Convective term. $\oint_{\partial\Omega_P} \mathbf{u}(\mathbf{u} \cdot \mathbf{n}) dS$ receives nothing from the wall, where $\mathbf{u} \cdot \mathbf{n}_w = 0$ for a wall at rest under either condition. On the wet face parts the mass flux is $\phi_f = \theta_f \bar{\mathbf{u}}_f \cdot \mathbf{S}_f$ and

Table 2: Volume-integrated terms on the fluid part of a cell for a body at rest; Γ is a diffusivity, p_f and $\bar{\mathbf{u}}_f$ are linear face interpolates, ϕ_f is the wet face flux.

term	wall contribution	discrete operator
unsteady	—	$\alpha_P V_P (\mathbf{u}_P^{n+1} - \mathbf{u}_P^n) / \Delta t$
convective	zero ($\mathbf{u} \cdot \mathbf{n}_w = 0$)	$\sum_f \phi_f \bar{\mathbf{u}}_f, \quad \phi_f = \theta_f \bar{\mathbf{u}}_f \cdot \mathbf{S}_f$
viscous, free slip	zero (no tangential traction)	$\sum_f \theta_f \Gamma_f \delta_f \mathbf{S}_f (\mathbf{u}_N - \mathbf{u}_P)$
viscous, no slip	$\nu (\mathbf{0} - \mathbf{u}_P) A_{\text{wall}} / d_{\text{wall}}$	same, plus $-\nu c_{\text{ns}} V_P \mathbf{u}_P$
pressure	$p_P \mathbf{S}_w$, absorbed by closure	$\sum_f \theta_f (p_f - p_P) \mathbf{S}_f$
continuity	zero	$\sum_f \phi_f = 0$

the transported value is the face interpolate (central differencing in all cases reported here, any OpenFOAM convection scheme in general).

Viscous term. $\oint \nu \nabla \mathbf{u} \cdot \mathbf{n} dS$ over the wet face parts is the standard two-point flux weighted by θ_f , with the diffusivity interpolated to the face and multiplied by θ_f ,

$$\sum_f \theta_f \nu_f |\mathbf{S}_f| \frac{\mathbf{u}_N - \mathbf{u}_P}{|\mathbf{d}_f|} \quad (+ \text{ the usual non-orthogonal correction on non-orthogonal meshes}). \quad (8)$$

On the wall segment, free slip means zero tangential traction and, since $\mathbf{u} \cdot \mathbf{n}_w = 0$, no normal viscous flux either, so the wall contributes nothing. For no slip the wall shear is approximated by a two-point difference between the cell value and the wall value $\mathbf{0}$ over the distance (6),

$$\int_{\text{wall}} \nu \nabla \mathbf{u} \cdot \mathbf{n} dS \simeq -\nu A_{\text{wall}} \frac{\mathbf{u}_P - \mathbf{0}}{d_{\text{wall}}} = -\nu c_{\text{ns}} V_P \mathbf{u}_P. \quad (9)$$

This is one implicit diagonal term per no-slip wall cell and per velocity component; the components stay uncoupled because the difference is taken for the full vector (the normal component is small anyway, being driven to zero by continuity).

Pressure term. The volume-integrated gradient is $\oint_{\partial\Omega_P} p \mathbf{n} dS$; splitting the boundary into the wet faces and the wall segment, and taking $p_w = p_P$ on the wall (first order),

$$\nabla_V p|_P \equiv \oint p \mathbf{n} dS = \sum_f \theta_f p_f \mathbf{S}_f + p_P \mathbf{S}_w \stackrel{\text{Prop. 1}}{=} \sum_f \theta_f (p_f - p_P) \mathbf{S}_f. \quad (10)$$

The wall pressure force — the form drag — is therefore absorbed exactly into the face sum. The momentum equation uses $\nabla_V p / V_P$, i.e. the gradient per unit *cell* volume; the factor α_P belongs to the mobility, not to the gradient (Section 5.2).

Continuity. $\oint \mathbf{u} \cdot \mathbf{n} dS = \sum_f \phi_f + \int_{\text{wall}} \mathbf{u} \cdot \mathbf{n}_w dS$, and the wall integral vanishes for a wall at rest. Hence $\sum_f \phi_f = 0$ with the same wet fluxes as in the convective term.

4.2 Properties

The following properties hold to round-off because every wall-related quantity is built from the same θ_f .

Proposition 2 (No force from a uniform pressure). *For $p \equiv p_0$, $\nabla_V p = \mathbf{0}$ in every cell, cut or not.*

Proof. $p_f - p_P = 0$ in (10). □

Proposition 3 (Free-stream divergence). *For a uniform discrete velocity $\mathbf{u} \equiv \mathbf{U}$, $\sum_f \phi_f = -\mathbf{U} \cdot \mathbf{S}_w$ in every cell. In particular the discrete divergence vanishes in a cut cell if and only if $\mathbf{U} \cdot \mathbf{n}_w = 0$.*

Proof. $\sum_f \phi_f = \mathbf{U} \cdot \sum_f \theta_f \mathbf{S}_f = -\mathbf{U} \cdot \mathbf{S}_w$ by (3). $\square$

Proposition 4 (No-penetration is implied, not imposed). *Let the velocity be locally uniform, $\mathbf{u} = \mathbf{U}$, on the faces of a wall cell with $A_{\text{wall}} \neq 0$. Discrete continuity $\sum_f \phi_f = 0$ holds if and only if $\mathbf{U} \cdot \mathbf{n}_w = 0$.*

Proof. Immediate from Proposition 3. The statement holds for any shape of the wall inside the cell, since only the area vector $\mathbf{S}_w$ enters. $\square$

Proposition 4 is the structural point of the method: the no-penetration condition is never written down, it is a consequence of enforcing mass conservation on the cut cell with fluxes and wall segment taken from the same fractions. The free-slip condition appears only as the *absence* of a wall term in the viscous operator, and the no-slip condition as the presence of the single diagonal term $\nu c_{\text{ns}} V_P$.

Proposition 5 (Small cells strengthen the diagonal). *In a cut cell with a planar wall segment, the no-slip coefficient (7) grows like $1/\alpha_P$ as $\alpha_P \rightarrow 0$ at fixed A_{wall} .*

Proof. $c_{\text{ns}} = 2A_{\text{wall}}^2/(\alpha_P V_P^2)$. $\square$

The small-cell limit therefore drives $\mathbf{u}_P$ onto the wall value and increases the diagonal dominance of the momentum rows instead of degrading them; together with sliver absorption it is what keeps the system well conditioned without cell merging. The price is paid by the pressure (Section 5.2, and Section 8).

Proposition 6 (Discrete conservation). *Summing the momentum equations of all cells, the convective and viscous face fluxes cancel in pairs, and the pressure term reduces to the pressure force on the bodies, $\sum_P \nabla_V p|_P = \sum_P p_P \mathbf{S}_{w,P} + (\text{domain boundary faces})$.*

Proof. Each internal face appears in the rows of its two cells with the same θ_f and opposite $\mathbf{S}_f$. In (10) write $\sum_f \theta_f p_f \mathbf{S}_f + p_P \mathbf{S}_w$; the face terms cancel pairwise. Solid cells have all faces closed and contribute nothing. $\square$

Consequently the force on the bodies is the reaction to exactly the wall terms assembled in the momentum equations (Section 5.5), and at a steady state in a periodic domain the total force equals the driving body force times the fluid volume. The same argument gives global mass conservation, $\sum_P \sum_f \phi_f = 0$, and it is what makes the pressure system compatible (Section 5.3).

Measured on manufactured solutions of Poisson and convection–diffusion problems on a disk, the free-slip (Neumann) treatment is second-order accurate and the no-slip (Dirichlet) treatment first-order accurate. The cause is the pair of two-point differences in (8) and (9): the cell value of a cut cell is the average over its fluid part and lives at the centroid of that part, not at the cell centre or at the distance d_{wall} from the wall.

5 Pressure–velocity coupling, forces and flow-rate forcing

5.1 Solid cells

Every face of a solid cell is closed, so its rows in the momentum and pressure matrices have no off-diagonal entries and are decoupled from the fluid. They receive regular diagonals so that the linear solvers see a well-posed system.

- *Momentum.* The blanking diagonal $\nu b_P V_P = \nu \sum_f |\mathbf{S}_f| \delta_f$, the viscous diagonal of an unobstructed cell. The right-hand side is zero in a solid cell (the gradient (10) vanishes, the forcing is weighted with $\alpha_P = 0$), so $\mathbf{u}_P = \mathbf{0}$ exactly.
- *Pressure.* The diagonal $-\bar{a}$, with $\bar{a}$ the mean magnitude of the diagonal over the non-solid cells, and a zero source, so $p_P = 0$.

Neither value affects the fluid. The coefficient of the momentum row matters only for the rate at which an iterative solver converges; its size makes the solid rows neither dominate nor spoil the scaling of the matrix. Velocity and pressure are also set to zero in the solid when the solver starts.

5.2 SIMPLEC with cut-cell operators

Write the assembled momentum equation of cell P as

$$a_P \mathbf{u}_P = \mathbf{H}_P(\mathbf{u}) - \nabla_V p|_P, \quad \mathbf{H}_P = \mathbf{b}_P - \sum_{nb} a_{nb} \mathbf{u}_{nb}, \quad (11)$$

with $a_{nb} \leq 0$ the off-diagonal coefficients (OpenFOAM's sign convention), $\mathbf{b}_P$ the explicit sources (old time level, wall and blanking sources, forcing) and $\nabla_V p$ the volume-integrated cut-cell gradient (10). With under-relaxation λ_U the diagonal is a_P/λ_U and $\mathbf{b}_P$ holds the relaxation source. Define, per unit volume as OpenFOAM does,

$$A_P = \frac{a_P}{V_P}, \quad H1_P = -\frac{1}{V_P} \sum_{nb} a_{nb}, \quad r_A = \frac{1}{A_P} = \frac{V_P}{a_P}, \quad r_{At} = \frac{1}{A_P - H1_P} = \frac{V_P}{a_P + \sum_{nb} a_{nb}}. \quad (12)$$

r_A is the mobility of SIMPLE and r_{At} that of SIMPLEC, the response of the cell velocity to a pressure gradient when the neighbours respond in the same way. Note that the momentum equation keeps the full cell volume in A : the fluid volume enters only through the unsteady diagonal $\alpha_P V_P/\Delta t$ and the wall term. The predictor is $\mathbf{H}/a_P$, and the Rhie–Chow flux of the predictor is

$$\phi_{H,f} = \theta_f (r_A \mathbf{H}/V)_f \cdot \mathbf{S}_f + \theta_f (r_{At} - r_A)_f |\mathbf{S}_f| \partial_n p_f^{\text{old}} \left[+ (r_A)_f \mathcal{D}_f \right], \quad (13)$$

with ∂_n the face-normal gradient of the **snGrad** scheme and $\mathcal{D}_f$ the transient Rhie–Chow correction (below). The pressure equation is the discrete divergence of the corrected flux, with the Laplacian weighted by the same θ_f as the flux,

$$\sum_f \theta_f (r_{At})_f |\mathbf{S}_f| \partial_n p_f = \sum_f \phi_{H,f}, \quad \phi_f = \phi_{H,f} - \theta_f (r_{At})_f |\mathbf{S}_f| \partial_n p_f, \quad (14)$$

and the cell velocity is corrected with the cut-cell gradient, $\mathbf{u}_P = r_A \mathbf{H}_P/V_P - (r_{At})_P \nabla_V p|_P/V_P + ((r_{At})_P - (r_A)_P) \nabla_V p^{\text{old}}|_P/V_P$ (the last two terms combine into the familiar SIMPLEC update with $\mathbf{H}/a_P$ modified by $(r_A - r_{At}) \nabla p^{\text{old}}$). Laplacian and flux in (14) carry the same θ_f , so a closed face transports neither flux nor pressure information, and ϕ_f is discretely divergence-free after the solve.

Floors of the mobility. In a cut cell dominated by its wall term, a_P can be close to $\sum_{nb} |a_{nb}|$ and the SIMPLEC denominator $A - H1$ tiny or negative. It is floored:

- transient solver: $A_P - H1_P \geq \alpha_P/\Delta t$, the unsteady diagonal of the cell's own fluid volume (zero in solid cells, where $A - H1 = A > 0$);
- steady solver: $A_P - H1_P \geq 10^{-3} A_P$, a pure safeguard.

r_A itself is not floored. The reference implementation of Section 7 instead floors a_P at $\alpha_P V_P/\Delta t$; the difference is invisible in the converged transient solutions but is one reason why the steady solver converges to a slightly different fixed point (the Rhie–Chow dissipation depends on the relaxation).

Transient Rhie–Chow correction. With `ddtCorr` active, the flux of the predictor carries the difference between the old flux and the interpolate of the old velocity,

$$\mathcal{D}_f = c_f \frac{\phi_f^n - \theta_f \bar{\mathbf{u}}_f^n \cdot \mathbf{S}_f}{\Delta t}, \quad c_f = 1 - \min\left(\frac{|\phi_f^n - \theta_f \bar{\mathbf{u}}_f^n \cdot \mathbf{S}_f|}{|\phi_f^n| + \epsilon}, 1\right), \quad (15)$$

the Euler form of OpenFOAM’s `ddtCorr` with θ -weighted interpolation. It is added to (13) with the coefficient $(r_A)_f$. Other time schemes fall back to the Euler form with a warning.

5.3 Compatibility and the pressure reference

The fluid region is a pure Neumann island for the pressure: no face of the fluid carries a pressure boundary condition in a periodic domain, and the solid rows are decoupled. The pressure Laplacian L restricted to the fluid cells is therefore singular with null vector $\mathbf{1}$ and $\mathbf{1}^\top L = \mathbf{0}$. The reference is imposed as in `fvMatrix::setReference`, by doubling one diagonal entry d_r (and adding $d_r p_{\text{ref}}$ to the source) rather than replacing a row.

Lemma 1 (Transparent reference). *Let $L' = L + d_r \mathbf{e}_r \mathbf{e}_r^\top$ and $L'p = \mathbf{b}$ (with $p_{\text{ref}} = 0$). Then $d_r p_r = \mathbf{1}^\top \mathbf{b}$. In particular, for a compatible right-hand side ($\mathbf{1}^\top \mathbf{b} = 0$) the reference is transparent, $p_r = 0$ and $Lp = \mathbf{b}$ exactly; any incompatibility is deposited at the reference cell as a point source.*

Proof. Left-multiply $L'p = \mathbf{b}$ by $\mathbf{1}^\top$ and use $\mathbf{1}^\top L = \mathbf{0}$. $\square$

The right-hand side of (14) is compatible because $\sum_P \sum_f \phi_{H,f}$ telescopes (each internal and coupled face appears twice with opposite sign). In a domain with inflow and outflow boundaries the steady solver additionally balances the boundary fluxes (`adjustPhi`). The reference cell must be a fluid cell; this is checked at start-up.

5.4 Algorithm

Algorithm 1 summarises one time step of the transient solver; the steady solver is the same loop without the unsteady term, with one pass per iteration and relaxation of $\mathbf{u}$ and p .

Algorithm 1: One time step of the transient cut-cell solver (`ibmPimpleFoam`).

```

1 for outer iteration  $k = 1, \dots, n_{\text{outer}}$  do
2   assemble  $\alpha V \partial_t \mathbf{u} + \nabla \cdot (\phi \mathbf{u}) - \nabla \cdot (\theta \nu \nabla \mathbf{u}) + \nu c_{\text{ns}} V \mathbf{u} + \nu b V \mathbf{u} = \mathbf{s}$ ;
3   solve the predictor with  $-\nabla_V p$  (optional);
4   for corrector  $j = 1, \dots, n_{\text{corr}}$  do
5      $r_A, r_{At}$  (floored),  $\mathbf{H}/a_P$ ,  $\phi_H$  by (13) with  $\theta_f$ ;
6     solve  $\sum_f \theta_f (r_{At})_f |\mathbf{S}_f| \partial_n p_f = \sum_f \phi_{H,f}$  with regular solid rows and the reference;
7      $\phi := \phi_H - \theta r_{At} |\mathbf{S}_f| \partial_n p$ ; correct  $\mathbf{u}$  with  $\nabla_V p$ ; update the flow-rate forcing;
```

5.5 Force on the body

By closure, the pressure force on a body is the sum of the wall terms of (10),

$$\mathbf{F}_p = \sum_{P \in \text{wall}} p_P \mathbf{S}_{w,P}, \quad \mathbf{F}_v = \sum_{P \in \text{wall}} \nu c_{\text{ns}} V_P (\mathbf{u}_P), \quad (16)$$

the viscous force being the reaction to the traction (9) assembled in the momentum equation (zero for free slip). Wall cells without a body are counted on the first body. The forces are per unit density; for a two-dimensional case they refer to the depth of the mesh, and are divided by it to obtain forces per unit depth. By Proposition 6, at a steady state in a periodic domain with forcing $\mathbf{g}$ the total force equals $\mathbf{g} \sum_P \alpha_P V_P$ exactly; this identity is used as a consistency check in Section 7.

5.6 Constant flow rate in periodic domains

A periodic domain is driven by a uniform body force $\mathbf{g} = g\mathbf{e}$ adjusted so that the bulk velocity $\bar{u} = \sum_P \alpha_P V_P \mathbf{e} \cdot \mathbf{u}_P / V_f$, $V_f = \sum_P \alpha_P V_P$, equals a target U . The force acts on the fluid only: the source of cell P is $g\mathbf{e} \alpha_P V_P$.

Gain. Adding a force increment δg changes the cell velocity by $\delta \mathbf{u}_P = \mathbf{e} \rho_P \alpha_P \delta g$, where ρ_P is the mobility of the cell to a force that is the same in the neighbours. For a uniform increment the neighbour couplings cancel in the momentum equation, so $\rho_P = 1/(A_P - H1_P) = r_{At}$, the SIMPLEC mobility. The bulk response is

$$\delta \bar{u} = \bar{r} \delta g, \quad \bar{r} = \frac{\sum_P (r_{At})_P \alpha_P^2 V_P}{V_f}, \quad \text{hence} \quad \delta g = \lambda \frac{U - \bar{u}}{\bar{r}}, \quad (17)$$

with a relaxation λ (1 by default); the extra factor α_P in $\bar{r}$ comes from the force being weighted with α_P and the average with $\alpha_P V_P$. After every momentum predictor and pressure correction, δg is computed, added to the accumulated increment, and the velocity is corrected by $\mathbf{e} (r_{At})_P \alpha_P \delta g$. The stock OpenFOAM controller uses the gain $1/A$, which underestimates the bulk response by the diffusion number $\nu \Delta t / h^2$ when the viscous term is implicit, and overshoots by two orders of magnitude in the cases of Section 7. The mobility is floored at $A - H1 \geq 10^{-3}A$. The target must be non-zero (the flow direction is taken from it).

6 Implementation in OpenFOAM

6.1 Explicit operators on an unmodified mesh

The mesh is left untouched; θ_f and α_P enter the standard operators explicitly. The momentum equation of the transient solver reads

```
fvm::ddt(ibm.alpha(), U) // alpha V/dt; alpha.oldTime() = alpha^n
+ fvm::div(phi, U) // phi = theta (U_f . Sf) = ibm.flux(U)
- fvm::laplacian(ibm.thetaInterpolate(nu), U) // theta-weighted, no wall term
+ fvm::Sp(nu*ibm.noSlipCoeff(), U) // wall traction, no-slip bodies
+ fvm::Sp(nu*ibm.blankCoeff(), U) // blanking of solid cells
== fvOptions(U) - ibm.grad(p) // sum_f theta_f (p_f - p_P) Sf / V
```

The pressure equation is

```
fvm::laplacian(ibm.thetaInterpolate(rAtU), p) == fvc::div(phiHbyA) // phiHbyA = ibm.flux(HbyA)
ibm.constrainPressure(pEqn); // regular rows for the solid cells
```

and `pEqn.flux()` is then automatically θ -weighted. Everything else is `simpleFoam` and `pimpleFoam`: the boundary constraints on $\mathbf{H}/a_P$ and p , the pressure reference, pressure relaxation, finite-volume options, residual control and the PIMPLE loops. The steady and transient solvers differ from their stock counterparts only in the lines above, the floors of Section 5.2 and the forces.

This choice keeps the full cell volume V_P in `fvMatrix::A()`, so that $r_A = V_P/a_P$ is exactly the mobility of (12). Scaling the face area vectors and cell volumes of the mesh in place would make all operators θ/α -aware implicitly, but would change the SIMPLEC mobility to $\alpha V_P/a_P$, break the interpolation weights on closed faces and require floors for solid cells; explicit operators are transparent and term-for-term identical to the formulation above.

6.2 Structure

`ibmBody` a body (Table 1): analytical, any closed `searchableSurface`, or a supplied fraction field, with its wall condition.

`cutCellGeometry` the fractions and derived fields of Section 3, the cut-cell operators (gradient, flux, θ -interpolation, `ddtCorr`), the solid-row regularisation, the forces.

`ibmMeanVelocityForce` the constant-flow-rate forcing of Section 5.6, as a finite-volume option.
`ibmSimpleFoam`, `ibmPimpleFoam` the solvers.
`ibmSetGeometry` computes and writes the geometry fields, and checks that the cut-cell gradient of a uniform field is zero (the closure identity).

The geometry computation is local except for the synchronisation of coupled faces and global reductions, so parallel runs need nothing else.

6.3 Input and output

The geometry is controlled by `constant/ibmProperties`:

```
alphaMin    0.02;    // sliver threshold
nSubSamples 32;      // n_s, samples per direction for alpha
nBisect     40;      // bisection steps per cut edge
writeFields true;    writeForces true;
bodies
{
    cylinder
    {
        type cylinder; centre (5 5 0); axis (0 0 1); radius 1;
        wallType noSlip;                // or freeSlip; invert true for fluid inside
    }
}
```

The flow-rate forcing is an `fvOptions` entry of type `ibmMeanVelocityForce` with the target `Ubar` and `relaxation`. The fields `ibmAlpha`, `ibmTheta`, `ibmSw`, `ibmWallDistance`, `ibmNoSlipCoeff` and `ibmBodyIndex` are written at start-up. The forces of Section 5.5 are printed every step and written to `postProcessing/ibmForces/<time>/forces.dat` (time, then the pressure and viscous force of each body).

7 Verification

Unless stated otherwise the cases use a circular cylinder of radius $R = 1$ centred in a 10×10 doubly periodic box discretised by 81×81 cells ($h = 0.1235$), $\nu = 1$, $\alpha_{\min} = 0.02$, $n_s = 32$ and central differencing for convection. The reference is an independent implementation of the same discretisation on a structured grid (direct sparse solves, exact pinning of the bulk velocity by superposition), so that the comparison isolates the OpenFOAM implementation from the formulation. It differs from the OpenFOAM implementation in details that do not change the discrete fixed point: it uses the level set $r^2 - R^2$, a penalty of 10^{12} instead of the blanking diagonal, floors a_P instead of $A - H1$, and pins the bulk velocity exactly instead of by the controller of Section 5.6. Forces are quoted per unit depth and density.

7.1 Body at rest

The flow is driven at bulk velocity 0.5 ($Re_D = 1$). Table 3 compares geometry, fields and drag. The fractions agree to 10^{-11} and the cell fractions exactly. The transient solver run with the reference’s pseudo-time step ($\Delta t = 65h^2/\nu$, Euler, one corrector, SIMPLEC, no `ddtCorr`) converges to the same discrete fixed point (velocity to 10^{-5} , drag to 2×10^{-6}); the steady solver differs through the relaxation dependence of the Rhie–Chow dissipation, the pressure differences being concentrated in the cut cells. The body given as an STL (a 3600-gon) or as a fraction field gives the accuracy of its geometry route. With two bodies of mixed wall conditions (a no-slip cylinder of radius 1 and a free-slip one of radius 0.7) the forces sum to the driving force times the fluid volume to all printed digits, $0.625915 + 0.202512 = 0.828427$, the check of Proposition 6;

the free-slip body carries no viscous force. A four-rank parallel run reproduces the serial drag to all printed digits.

Table 3: Body at rest, no slip unless stated; reference drag 5.660008 (no slip), 2.935529 (free slip).

	steady	transient, ref. Δt	STL body	fraction field	free slip, transient
α , cell by cell	exact	exact	exact	exact	exact
θ , max. diff.	7.6×10^{-11}	7.6×10^{-11}	6.4×10^{-6}	0.11	7.6×10^{-11}
d_{wall} , max. diff.	10^{-11}	10^{-11}	2.8×10^{-6}	2.2×10^{-2}	10^{-11}
u , max. diff. / max u	2.2×10^{-4}	9.4×10^{-6}	2.2×10^{-4}	3.3×10^{-3}	1.7×10^{-5}
p , max. diff. / max p	3.0×10^{-2}	9.6×10^{-4}	3.0×10^{-2}	9.8×10^{-2}	1.9×10^{-4}
drag	5.660383	5.659997	5.660383	5.651222	2.935496




Figure 4: Body at rest, transient solver with the reference’s pseudo-time step: streamwise velocity (top) and pressure (bottom) against the reference, difference field, and profiles through the cylinder centre.

8 Limitations

- *Order.* First order in time (implicit Euler). First order in space for no-slip and fixed-value walls, second order for free slip and zero flux; the cut-cell pressure and the forces are first order.
- *Cut-cell pressure.* The mobility of a cut cell scales with its fluid fraction, so its pressure is poorly determined; the wall pressure is taken as the cell pressure. Cell merging or a wall pressure from the full cells is the remedy.
- *Slivers.* Absorption below $\alpha_{\min}$ displaces the boundary by $O(\alpha_{\min}h)$ and removes the fluid volume of the absorbed cell. Mass conservation of the flow is unaffected. Cell merging is the clean remedy. A cut cell whose faces all close (thin features, or after coupled-face synchronisation) while its α stays above $\alpha_{\min}$ is not guarded against. $\alpha_{\min}$ must be positive, otherwise cells with $\alpha = 0$ from a fraction field keep open faces.

- *Sub-cell features.* Faces are classified by their vertices and cells by their vertices, face centres and centre; a feature that passes between them is not seen.
- *Laminar.* Turbulence models are constructed for their effective viscosity only; their own transport equations would need the same treatment as a passive scalar. `fixedFluxPressure` boundary conditions are not θ -aware.
- *Fraction-field bodies.* Reconstructing θ from the 0.5 iso-surface of the point-interpolated fraction smears the wall over one cell (drag -0.16% on the cylinder, θ errors up to 0.11); a per-cell iso-value reproducing each cell's α would bring it to $O(h^2)$.
- *One fluid region.* A single pressure reference; a body splitting the fluid would need one per region.

9 Conclusions

The method rests on one geometric statement, the closure identity of a cut cell. From the closure identity follow exactness for uniform flow, the absence of spurious pressure forces, no-penetration itself, discrete conservation of momentum and exact compatibility of the pressure system. The whole boundary treatment adds to the interior scheme two scalar fields, θ_f and α_P , one diagonal coefficient for a no-slip wall. Implemented in OpenFOAM with explicit operators on an unmodified mesh, it reproduces an independent reference implementation to the level of the linear-solver tolerances for a body at rest.